

# Build a 12 Factor Microservice in an Hour

Andrew Rouse: Open Liberty Developer for MicroProfile and CDI, IBM

@azquelt

# Contents

Basic concepts of 12 factor apps

Demo of creating a 12 factor microservice using MicroProfile

# 12 Factors in a nut shell




- A methodology
- Best Practices
- Manifesto

<https://12factor.net/> by Heroku

# Why 12 factor?

- Define the contract between applications and infrastructure



# THE FACTORS

1. Codebase
2. Dependencies
3. Config
4. Backing Services
5. Build, Release, Run
6. Processes
7. Port binding
8. Concurrency
9. Disposability
10. Dev / Prod parity
11. Logs
12. Admin Processes

# Why 12 factor?

In the modern era, software is commonly delivered as a service: called *web apps*, or *software-as-a-service*. The twelve-factor app is a methodology for building software-as-a-service apps that:

Use **declarative** formats for setup automation, to minimize time and cost for new developers joining the project;

Have a **clean contract** with the underlying operating system, offering **maximum portability** between execution environments;

Are suitable for **deployment** on modern **cloud platforms**, obviating the need for servers and systems administration;

**Minimize divergence** between development and production, enabling **continuous deployment** for maximum agility;


And can **scale up** without significant changes to tooling, architecture, or development practices.

The twelve-factor methodology can be applied to apps written in any programming language, and which use any combination of backing services (database, queue, memory cache, etc).

# I. Codebase

“One codebase tracked in revision control, many deploys.”

- Dedicate smaller teams to individual applications or microservices.
  - Following the discipline of single repository for an application forces the teams to analyze the seams of their application, and identify potential monoliths that should be split off into microservices.
- Use a single source code repository for a single application (1:1 relation).
- E.g. use a central git repo




# II. Dependencies

“Explicitly declare and isolate dependencies”

A cloud-native application does not rely on the pre-existence of dependencies in a deployment target.

Developer Tools declare and isolate dependencies

- [Maven](#) and [Gradle](#) declare Java library dependencies
- Dockerfile declares dependencies on OS and other tools
- Each microservice declares its own dependencies






# III. Config



“Store config in the environment”


- Changing config should not need to repackaging your application
- Use Kubernetes configmaps and secrets for container services, rather than values specified in the app or container image
- Use MicroProfile Config to inject the config properties into the microservices



## IV. Backing services



“Treat backing services as attached resources”



# V. Build, release, run

“Strictly separate build and run stages”

- Source code is used in the build stage. Configuration data is added to define a release stage that can be deployed. Any changes in code or config will result in a new build/release
- Needs to be considered in CI pipeline

## **IBM**

- [UrbanCode Deploy](#)
- [IBM Cloud Continuous Delivery Service](#)

## **AWS**

- [AWS CodeBuild](#)
- [AWS CodeDeploy](#)
- [AWS CodePipeline](#) (not yet integrated with EKS)

## **Azure**

- [Visual Studio Team Services \(VSTS\) \(includes git\)](#)
- [Web App for Containers](#)  
feature of Azure App Service

# VI. Processes

“Execute the app as one or more stateless processes”

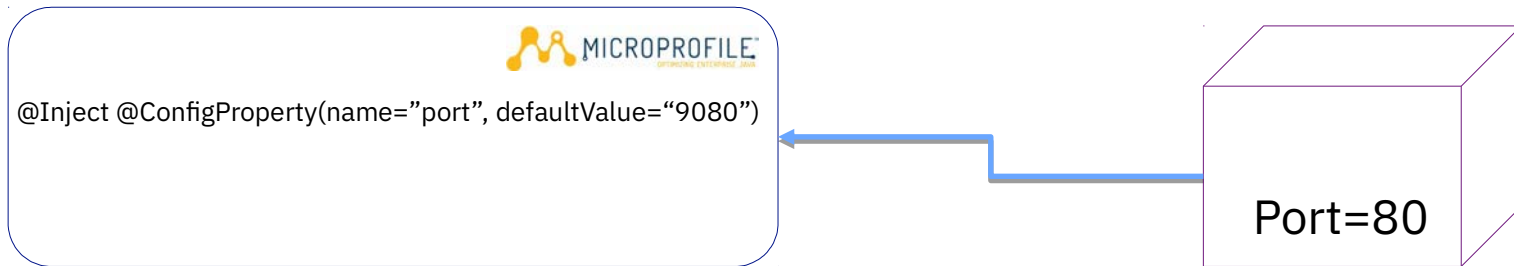
- Stateless and share-nothing
- Any persistence uses a backing service
  - Database
  - Shared cache
  - Message queue

# VII. Port binding



“Export services via port binding”


- Applications are fully self-contained and expose services only through ports. Port assignment is done by the execution environment
- Ingress/service definition of k8s manages mapping of ports
- Use MP Config to inject ports to microservices to connect up microservices



# VIII. Concurrency

“Scale out via the process model”

- Applications use processes independent from each other to scale out (allowing for load balancing)
- To be considered in application design
- Cloud autoscaling services: [auto]scaling built into k8s
- Build microservices



# IX. Disposability



“Maximize robustness with fast startup and graceful shutdown”

- Processes start up fast.
- Processes shut down gracefully when requested.
- Processes are robust against sudden death
  - Ensure you won't lose data if part of your system dies
  - Use MicroProfile Fault Tolerance to make calls to other services resilient

# X. Dev/prod parity

“Keep development, staging, and production as similar as possible”

- Development and production are as close as possible (in terms of code, people, and environments)
- Can use helm to deploy in repeatable manner
- Use (name)spaces for isolation of similar setups



# XI. Logs

“Treat logs as event streams”




- App writes all logs to stdout
- Use a structured output for meaningful logs suitable for analysis. Execution environment handles routing and analysis infrastructure

## XII. Admin processes

“Run admin/management tasks as one-off processes”

- Tooling: standard k8s tooling like “kubectl exec” or Kubernetes Jobs
- Also to be considered in solution/application design
- For example, if an application needs to migrate data into a database, place this task into a separate component instead of adding it to the main application code at startup

# THE FACTORS

1. Codebase
2. Dependencies
3. Config  MICROPROFILE™  
OPTIMIZING ENTERPRISE JAVA
4. Backing Services
5. Build, Release, Run
6. Processes
7. Port binding  MICROPROFILE™  
OPTIMIZING ENTERPRISE JAVA
8. Concurrency
9. Disposability  MICROPROFILE™  
OPTIMIZING ENTERPRISE JAVA
10. Dev / Prod parity
11. Logs
12. Admin Processes

# DEMO


|               |           |             |                         |                              |                |                      |              |                    |  |            |                                       |
|---------------|-----------|-------------|-------------------------|------------------------------|----------------|----------------------|--------------|--------------------|--|------------|---------------------------------------|
| 1<br>Codebase | 2<br>Deps | 3<br>Config | 4<br>Backing<br>Service | 5<br>Build<br>Release<br>Run | 6<br>Stateless | 7<br>Port<br>Binding | 8<br>Scaling | 9<br>Disposability | <del>10<br/>Prod, Dev,<br/>Staging</del> | 11<br>Logs | <del>12<br/>Admin<br/>Processes</del> |
|---------------|-----------|-------------|-------------------------|------------------------------|----------------|----------------------|--------------|--------------------|--|------------|---------------------------------------|

1. Create a REST app (stateless, in one codebase)
2. Declare our dependencies using Maven and Dockerfile
  - Including OpenLiberty
3. Release and Run by deploying with config to Kubernetes
4. Update our app to call another service
5. Deploy again with scaling

**12 FACTOR BINGO**


# References

- Code sample to demonstrate 12-factor app
  - <https://github.com/Azquelt/12factor-deployment>
  - <https://github.com/Azquelt/12factor-app-a>
  - <https://github.com/Azquelt/12factor-app-b>
- <http://microprofile.io>
- <http://openliberty.io> (especially guides)
- <https://www.12factor.net/>



# 12 factor app

- Use MicroProfile and K8s to build a microservice => 12 factor app



# MicroProfile Config



## Why?

- Configure Microservice without repacking the application

– Access configuration via

- Programmatic lookup

```
Config config = ConfigProvider.getConfig();  
config.getValue("myProp", String.class);
```

## How?

- Specify the configuration in configuration sources

- CDI Injection

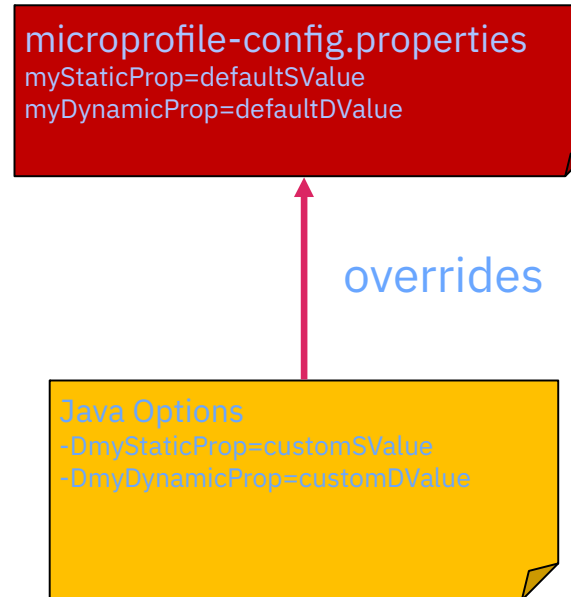
```
@Inject  
@ConfigProperty(name="my.string.property")  
String myProp;
```

## Static Config

```
@Inject  
@ConfigProperty(name="myStaticProp")  
private String staticProp;
```

## Dynamic Config

```
@Inject  
@ConfigProperty(name="myDynamicProp")  
private Provider<String> dynamicProp;
```





# MicroProfile Fault Tolerance

A solution to build a resilient microservice

- ❖ Retry - `@Retry`
- ❖ Circuit Breaker - `@CircuitBreaker`
- ❖ Bulk Head - `@Bulkhead`
- ❖ Time out - `@Timeout`
- ❖ Fallback - `@Fallback`

# Backup: Using IBM Cloud Private

|                            |   |
|----------------------------|---|
| <b>Codebase</b>            | Source: Github Enterprise, github<br>Images: any registry, IBM Cloud private registry                           |
| <b>Dependencies</b>        | Dependency management of language environment; container build process for repeatable inclusion of dependencies |
| <b>Config</b>              | k8s configmaps and secrets  |
| <b>Backing services</b>    | Use configuration (see previous factor) to define target server as used by application                          |
| <b>Build, release, run</b> | UrbanCode Deploy<br>UrbanCode Release<br>Plus k8s mechanisms with CI tooling                                    |
| <b>Processes</b>           | To be considered in application design  |

|                        |  |
|------------------------|--|
| <b>Port binding</b>    | Application needs to expose ports.<br>Ingress/service definition of k8s manages mapping of ports |
| <b>Concurrency</b>     | App design ([auto]scaling built into k8s)  |
| <b>Disposability</b>   | App design   |
| <b>Dev/prod parity</b> | Can use helm to deploy in same way. Namespaces for isolation of similar areas                    |
| <b>Logs</b>            | ELK as part of ICP (or RYO)  |
| <b>Admin processes</b> | App design; standard k8s tooling like “kubectl exec” or Kubernetes Jobs                          |

# Build a 12 Factor Microservice in an Hour

Andrew Rouse: Open Liberty Developer for MicroProfile and CDI, IBM

@azquelt

## Contents

Basic concepts of 12 factor apps

Demo of creating a 12 factor microservice using MicroProfile

# 12 Factors in a nut shell



- A methodology
- Best Practices
- Manifesto

<https://12factor.net/> by Heroku

## Why 12 factor?

- Define the contract between applications and infrastructure



- Application design influences the infrastructure required (Kubernetes, Database selection, etc...)
- Application design guides price; a cloud-native application can be resilient without surplus infrastructure
- Some services are provided by an application but some are provided from the infrastructure
- Applications depend on the features and services of infrastructure to support agile development.
- Infrastructure requires applications to expose endpoints and integrations to be managed autonomously (e.g. Kubernetes asks applications to expose health endpoint)
- Increased resource and memory consumption of independently running components that

## THE FACTORS

- |                        |                      |
|------------------------|----------------------|
| 1. Codebase            | 7. Port binding      |
| 2. Dependencies        | 8. Concurrency       |
| 3. Config              | 9. Disposability     |
| 4. Backing Services    | 10.Dev / Prod parity |
| 5. Build, Release, Run | 11.Logs              |
| 6. Processes           | 12.Admin Processes   |

## Why 12 factor?

In the modern era, software is commonly delivered as a service: called *web apps*, or *software-as-a-service*. The twelve-factor app is a methodology for building software-as-a-service apps that:

Use **declarative** formats for setup automation, to minimize time and cost for new developers joining the project;

Have a **clean contract** with the underlying operating system, offering **maximum portability** between execution environments;

Are suitable for **deployment** on modern **cloud platforms**, obviating the need for servers and systems administration;

**Minimize divergence** between development and production, enabling **continuous deployment** for maximum agility;

And can **scale up** without significant changes to tooling, architecture, or development practices.

The twelve-factor methodology can be applied to apps written in any programming language, and which use any combination of backing services (database, queue, memory cache, etc).

From <https://12factor.net>




## I. Codebase

“One codebase tracked in revision control, many deploys.”

- Dedicate smaller teams to individual applications or microservices.
- Following the discipline of single repository for an application forces the teams to analyze the seams of their application, and identify potential monoliths that should be split off into microservices.

➤ Use a single source code repository for a single application (1:1 relation).  
➤ E.g. use a central git repo



A *codebase* is any single repo (in a centralized revision control system like Subversion), or any set of repos who share a root commit (in a decentralized system like Git where changes are committed.) Also in Git – scenario of repository per package/service.

### multiple repository advantages:

- Clear ownership: team that owns a service is clearly responsible for independently develop and deploy the full stack of that service
- Smaller code base: Separate repositories for a service leads to smaller code base and lesser complexity during code merge.
- Narrow clones: faster DevOps and automated build and releases as smaller code base means less code download/clone time

### multiple repository disadvantages:

**Difficult development and debugging:** development, cross team communications and shared codes are difficult to maintain and thus development and debugging can be an issue

**Abstracts the knowledge of the platform:** Since each team is only responsible for a single service, integration becomes an issue and knowledge of the platform can decrease

**Multiple Everything:** Multiple dependencies, duplication and integrations


## II. Dependencies

## “Explicitly declare and isolate dependencies”

A cloud-native application does not rely on the pre-existence of dependencies in a deployment target.

## Developer Tools declare and isolate dependencies

- [Maven](#) and [Gradle](#) declare Java library dependencies
- Dockerfile declares dependencies on OS and other tools
- Each microservice declares its own dependencies



### III. Config



“Store config in the environment”

- Changing config should not need to repackage your application
- Use Kubernetes configmaps and secrets for container services, rather than values specified in the app or container image
- Use MicroProfile Config to inject the config properties into the microservices




*Do these make sense? –HJ  
(struggling how these offerings help pushing  
configurations –ub)*

Let's leave them because we have the space.  
I'm not really attached to them, so we can  
delete...

## IV. Backing services



“Treat backing services as attached resources”



- A backing service is any service on which an application relies (data stores, messaging systems, caching systems, security services)
- Use configuration (see factor III) to define target service as used by application. Configuration defines access to a backing service
- **Resource Binding** should be done via external configuration.
- Attach and detach backing services from an application at will, without re-deploying the application.
- Fault Tolerance pattern: allow code to stop communicating with misbehaving backing services, providing a fallback path

## V. Build, release, run

“Strictly separate build and run stages”

- Source code is used in the build stage. Configuration data is added to define a release stage that can be deployed. Any changes in code or config will result in a new build/release
- Needs to be considered in CI pipeline

### IBM

- [UrbanCode Deploy](#)
- [IBM Cloud Continuous Delivery Service](#)

### AWS

- [AWS CodeBuild](#)
- [AWS CodeDeploy](#)
- [AWS CodePipeline](#) (not yet integrated with EKS)

### Azure

- [Visual Studio Team Services \(VSTS\) \(includes git\)](#)
- [Web App for Containers](#) feature of Azure App Service

## VI. Processes

“Execute the app as one or more stateless processes”

- Stateless and share-nothing
- Any persistence uses a backing service
  - Database
  - Shared cache
  - Message queue


- Components are stateless and shared-nothing. State can be put in a stateful backing service (database)
- Stateless components can be replaced quickly if they fail
- Avoid dependencies on sticky sessions and keep session data in a persistent store to ensure traffic can be routed to other processes without service disruption
- To be considered in application design by the developer, not in Cloud roll-out

## VII. Port binding



“Export services via port binding”


- Applications are fully self-contained and expose services only through ports. Port assignment is done by the execution environment
- Ingress/service definition of k8s manages mapping of ports
- Use MP Config to inject ports to microservices to connect up microservices



## VIII. Concurrency

“Scale out via the process model”

- Applications use processes independent from each other to scale out (allowing for load balancing)
- To be considered in application design
- Cloud autoscaling services: [auto]scaling built into k8s
- Build microservices





## IX. Disposability



“Maximize robustness with fast startup and graceful shutdown”

- Processes start up fast.
- Processes shut down gracefully when requested.
- Processes are robust against sudden death
  - Ensure you won't lose data if part of your system dies
  - Use MicroProfile Fault Tolerance to make calls to other services resilient

## X. Dev/prod parity

“Keep development, staging, and production as similar as possible”

- Development and production are as close as possible (in terms of code, people, and environments)
- Can use helm to deploy in repeatable manner
- Use (name)spaces for isolation of similar setups

## XI. Logs

“Treat logs as event streams”




- App writes all logs to stdout
- Use a structured output for meaningful logs suitable for analysis. Execution environment handles routing and analysis infrastructure

## XII. Admin processes

“Run admin/management tasks as one-off processes”

- Tooling: standard k8s tooling like “kubectl exec” or Kubernetes Jobs
- Also to be considered in solution/application design
- For example, if an application needs to migrate data into a database, place this task into a separate component instead of adding it to the main application code at startup

## THE FACTORS

1. Codebase
2. Dependencies
3. Config  MICROPROFILE
4. Backing Services
5. Build, Release, Run
6. Processes
7. Port binding  MICROPROFILE
8. Concurrency
9. Disposability  MICROPROFILE
- 10.Dev / Prod parity
- 11.Logs
- 12.Admin Processes

## DEMO



1. Create a REST app (stateless, in one codebase)
2. Declare our dependencies using Maven and Dockerfile
  - Including OpenLiberty
3. Release and Run by deploying with config to Kubernetes
4. Update our app to call another service
5. Deploy again with scaling

Not covering 10 or 12 because...

We will create a REST app which is \*stateless and has \*one codebase

Declare our \*dependencies using maven and a Dockerfile

Includes OpenLiberty which will take care of \*binding to a port, starting up quickly and shutting down gracefully which is important for \*disposability, and will \*log to stdout

We will then build our app before doing a \*release and run by deploying the app with its config to kubernetes


Then we'll update our app to call another service which we'll treat as an \*attached backing service

And finally we'll deploy it again and demonstrate \*scaling with Kubernetes

Hopefully allowing us to achieve \*12 Factor Bingo!

## References


- Code sample to demonstrate 12-factor app
  - <https://github.com/Azquelt/12factor-deployment>
  - <https://github.com/Azquelt/12factor-app-a>
  - <https://github.com/Azquelt/12factor-app-b>
- <http://microprofile.io>
- <http://openliberty.io> (especially guides)
- <https://www.12factor.net/>



- Application design influences the infrastructure required (Kubernetes, Database selection, etc...)
- Application design guides price; a cloud-native application can be resilient without surplus infrastructure
- Some services are provided by an application but some are provided from the infrastructure
- Applications depend on the features and services of infrastructure to support agile development.
- Infrastructure requires applications to expose endpoints and integrations to be managed autonomously (e.g. Kubernetes asks applications to expose health endpoint)
- Increased resource and memory consumption of independently running components that

## 12 factor app

- Use MicroProfile and K8s to build a microservice => 12 factor app



- Application design influences the infrastructure required (Kubernetes, Database selection, etc...)
- Application design guides price; a cloud-native application can be resilient without surplus infrastructure
- Some services are provided by an application but some are provided from the infrastructure
- Applications depend on the features and services of infrastructure to support agile development.
- Infrastructure requires applications to expose endpoints and integrations to be managed autonomously (e.g. Kubernetes asks applications to expose health endpoint)
- Increased resource and memory consumption of independently running components that



# MicroProfile Config



## Why?

- Configure Microservice without repacking the application

– Access configuration via

- Programmatic lookup

```
Config config = ConfigProvider.getConfig();  
config.getValue("myProp", String.class);
```

## How?

- Specify the configuration in configuration sources

- CDI Injection

```
@Inject  
@ConfigProperty(name="my.string.property")  
String myProp;
```

## Static Config

```
@Inject  
@ConfigProperty(name="myStaticProp")  
private String staticProp;
```

## Dynamic Config

```
@Inject  
@ConfigProperty(name="myDynamicProp")  
private Provider<String> dynamicProp;
```

microprofile-config.properties  
myStaticProp=defaultSValue  
myDynamicProp=defaultDValue

overrides

Java Options  
-DmyStaticProp=customSValue  
-DmyDynamicProp=customDValue

# MicroProfile Fault Tolerance

A solution to build a resilient microservice

- ❖ Retry - `@Retry`
- ❖ Circuit Breaker - `@CircuitBreaker`
- ❖ Bulk Head - `@Bulkhead`
- ❖ Time out - `@Timeout`
- ❖ Fallback - `@Fallback`

# Backup: Using IBM Cloud Private

|                            |   |
|----------------------------|---|
| <b>Codebase</b>            | Source: Github Enterprise, github<br>Images: any registry, IBM Cloud private registry                           |
| <b>Dependencies</b>        | Dependency management of language environment; container build process for repeatable inclusion of dependencies |
| <b>Config</b>              | k8s configmaps and secrets  |
| <b>Backing services</b>    | Use configuration (see previous factor) to define target server as used by application                          |
| <b>Build, release, run</b> | UrbanCode Deploy<br>UrbanCode Release<br>Plus k8s mechanisms with CI tooling                                    |
| <b>Processes</b>           | To be considered in application design  |

|                        |  |
|------------------------|--|
| <b>Port binding</b>    | Application needs to expose ports.<br>Ingress/service definition of k8s manages mapping of ports |
| <b>Concurrency</b>     | App design ([auto]scaling built into k8s)  |
| <b>Disposability</b>   | App design   |
| <b>Dev/prod parity</b> | Can use helm to deploy in same way. Namespaces for isolation of similar areas                    |
| <b>Logs</b>            | ELK as part of ICP (or RYO)  |
| <b>Admin processes</b> | App design; standard k8s tooling like "kubectl exec" or Kubernetes Jobs                          |